

# Linear Classifiers, Logistic Regression, Multiclass Classification

## Overview

**Classification:** predicting a discrete-valued target

- ▶ **Binary classification:** predicting a binary-valued target
- ▶ **Multiclass classification:** predicting a discrete( $> 2$ )-valued target

## Binary linear classification

**classification:** given a  $D$ -dimensional input  $\mathbf{x} \in \mathbb{R}^D$  predict a discrete-valued target

**binary:** predict a binary target  $t \in \{0, 1\}$

- ▶ Training examples with  $t = 1$  are called **positive examples**, and training examples with  $t = 0$  are called **negative examples**. Sorry.
- ▶  $t \in \{0, 1\}$  or  $t \in \{-1, +1\}$  is for computational convenience.

**linear:** model prediction  $y$  is a linear function of  $\mathbf{x}$ , followed by a threshold  $r$ :

$$z = \mathbf{w}^\top \mathbf{x} + b$$
$$y = \begin{cases} 1 & \text{if } z \geq r \\ 0 & \text{if } z < r \end{cases}$$

## Simplifications

我们对上面的 binary linear classification 进行简化。

- Eliminating the threshold (消除阈值)

我们假设 threshold  $r = 0$  (该假设“不失一般性”, without loss of generality or WLOG, 即该个例能代表普遍情况, 而非一种特例):

$$\mathbf{w}^\top \mathbf{x} + b \geq r \iff \mathbf{w}^\top \mathbf{x} + \underbrace{b - r}_{\triangleq w_0} \geq 0$$

注:  $\triangleq$  为恒等式, 即  $b - r = w_0$

- Eliminating the bias

在上式中， $b$  是 bias。所以，可以和 linear regression 中一样：

在  $w$  中增加一列  $w_0$ ，使  $w_0 = b$ （实际上是  $b - r$ ），然后增加一个 dummy feature  $x_0$ （ $x_0$  永远为 1）。这样在运算的时候是  $w_0^\top x_0 = b$ 。

- Simplified model


Receive input  $\mathbf{x} \in \mathbb{R}^{D+1}$  with  $x_0 = 1$ :

$$z = \mathbf{w}^\top \mathbf{x}$$

$$y = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## The Geometric Picture

### Input Space, or Data Space for NOT example



注：NOT example 是指  $t$  的结果为 NOT  $x_1$ 。


上图中蓝色的线为 decision boundary（在 2D 中，它是一条线；在高维中，它是一个 hyperplane）： $\{\mathbf{x} : \mathbf{w}^\top \mathbf{x} = 0\}$

decision boundary 划分出了两个 half-spaces:

$H_+ = \{\mathbf{x} : \mathbf{w}^\top \mathbf{x} \geq 0\}$ ,  $H_- = \{\mathbf{x} : \mathbf{w}^\top \mathbf{x} < 0\}$ 。图中红绿色的加减号所在的位置，对应右边表格中  $x_0$  和  $x_1$  的值。上方的红色区域是 negative space，下方的绿色区域是 positive space（对于  $t$ ）。

如果这个 boundary 可以完美地将 training examples 区分开，我们就说 data is linearly separable。

## Weight Space





根据上表可得：

- ▶ When  $x_1 = 0$ , need:  $z = w_0 x_0 + w_1 x_1 \geq 0 \iff w_0 \geq 0$
- ▶ When  $x_1 = 1$ , need:  $z = w_0 x_0 + w_1 x_1 < 0 \iff w_0 + w_1 < 0$

因此，蓝色叉号所在的区域就是  $w_0$  和  $w_1$  可以取值的区域，叫做 feasible region。

## Visualizations of the AND example

$x_0$	$x_1$	$x_2$	$t$
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	1



- Slice for  $x_0 = 1$  and
- example sol:  $w_0 = -1.5$ ,  $w_1 = 1$ ,  $w_2 = 1$
- decision boundary:  
 $w_0 x_0 + w_1 x_1 + w_2 x_2 = 0$   
 $\implies -1.5 + x_1 + x_2 = 0$

- Slice for  $w_0 = -1.5$  for the constraints
- $w_0 < 0$
- $w_0 + w_2 < 0$
- $w_0 + w_1 < 0$
- $w_0 + w_1 + w_2 \geq 0$

## Towards Logistic Regression

### Loss Function

Seemingly obvious loss function: 0-1 loss

$$\begin{aligned}\mathcal{L}_{0-1}(y, t) &= \begin{cases} 0 & \text{if } y = t \\ 1 & \text{if } y \neq t \end{cases} \\ &= \mathbb{I}[y \neq t]\end{aligned}$$

Usually, the cost  $\mathcal{J}$  is the averaged loss over training examples; for 0 - 1 loss, this is the misclassification rate (错分类率):

$$\mathcal{J} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}[y^{(i)} \neq t^{(i)}]$$

接下来，要对模型进行优化。

### Attempt 1: 0-1 loss


Minimum of a function will be at its critical points.

Let's try to find the critical point of 0-1 loss

Chain rule:

$$\frac{\partial \mathcal{L}_{0-1}}{\partial w_j} = \frac{\partial \mathcal{L}_{0-1}}{\partial z} \frac{\partial z}{\partial w_j}$$

But  $\partial \mathcal{L}_{0-1} / \partial z$  is zero everywhere it's defined!



- ▶  $\partial \mathcal{L}_{0-1} / \partial w_j = 0$  means that changing the weights by a very small amount probably has no effect on the loss.
- ▶ Almost any point has 0 gradient!

关于 0-1 loss，还有一个问题是：根据最终预测来定义，这本质上是不连续的。所以 0-1 loss 不能用于优化模型。


## Attempt 2: Linear Regression

有时，我们可以用更容易优化的 loss function 来替换当前的 loss function。这被称为 relaxation with a smooth surrogate loss function（用光滑的替代函数来放松）。

$$z = \mathbf{w}^\top \mathbf{x}$$

$$\mathcal{L}_{\text{SE}}(z, t) = \frac{1}{2}(z - t)^2$$

The problem:



对于该 loss function,  $z$  的阈值为  $\frac{1}{2}$  时模型达到最优 (看上图, 如果阈值为其他数, 拟合的线可能是斜的, 这样会增大 loss, 只有拟合线为  $t = \frac{1}{2}$  时, loss 最小)。

由于上述原因, 使用该 loss function 不能很好的进行优化。


### Attempt 3: Logistic Activation Function

现在我们将  $y$  压缩到区间  $[0, 1]$  之间。

The **logistic function** is a kind of **sigmoid**, or S-shaped function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$\sigma^{-1}(y) = \log(y/(1 - y))$  is called the **logit**.



注: logit 和 sigmoid 互为反函数。

具有逻辑非线性 (logistic nonlinearity) 的线性模型称为: log-linear:

$$z = \mathbf{w}^\top \mathbf{x}$$


$$y = \sigma(z)$$

$$\mathcal{L}_{SE}(y, t) = \frac{1}{2}(y - t)^2.$$

Used in this way,  $\sigma$  is called an **activation function**.

### The problem:

(plot of  $\mathcal{L}_{SE}$  as a function of  $z$ , assuming  $t = 1$ )



$$\frac{\partial \mathcal{L}}{\partial w_j} = \frac{\partial \mathcal{L}}{\partial z} \frac{\partial z}{\partial w_j}$$


- 由前面  $\sigma(z)$  的函数图像可知: For  $z \ll 0$ ,  $\sigma(z) \approx 0$ .
- 所以  $\frac{\partial \mathcal{L}}{\partial z} \approx 0 \implies \frac{\partial \mathcal{L}}{\partial w_j} \approx 0 \implies$  derivative w.r.t. (with respect to)  $w_j$  is small  $\implies w_j$  is like a critical point
- 因此, 用该 function, 可能使模型向负轴方向优化, 从而不能到达真正的 critical point

# Logistic Regression

## Cross-entropy loss

Cross-entropy loss (aka log loss) captures this intuition:

$$\begin{aligned}\mathcal{L}_{\text{CE}}(y, t) &= \begin{cases} -\log y & \text{if } t = 1 \\ -\log(1 - y) & \text{if } t = 0 \end{cases} \\ &= -t \log y - (1 - t) \log(1 - y)\end{aligned}$$



注：交叉熵 loss 的机制是这样的：如果对于一个 sample，预测的是 1（一般还会包含一个置信度，比如有 90% 的把握认为是 1），但实际结果是 0，这时 cross-entropy loss 会对这个结果进行惩罚（一般，错误预测的置信度越大，惩罚越大），即让 loss 变得很大（见上图蓝线）。


## Logistic Regression

$$z = \mathbf{w}^T \mathbf{x}$$

$$y = \sigma(z)$$

$$= \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}_{\text{CE}} = -t \log y - (1 - t) \log(1 - y)$$



Plot is for target  $t = 1$ .

## Gradient Descent for Logistic Regression

由于 logistic loss 是一个凸函数（**convex function**，如上图，曲线上方空间是凸出来的函数），所以我们可以用 gradient descent 来对其进行优化。

## Gradient of Logistic Loss

我们首先将模型的权重进行合理的初始化，这里权重为  $\mathbf{w}$ ，我们将  $\mathbf{w}$  初始化为 0。

Back to logistic regression:

$$\mathcal{L}_{\text{CE}}(y, t) = -t \log(y) - (1 - t) \log(1 - y)$$
$$y = 1/(1 + e^{-z}) \quad \text{and} \quad z = \mathbf{w}^\top \mathbf{x}$$

Therefore

$$\frac{\partial \mathcal{L}_{\text{CE}}}{\partial w_j} = \frac{\partial \mathcal{L}_{\text{CE}}}{\partial y} \cdot \frac{\partial y}{\partial z} \cdot \frac{\partial z}{\partial w_j} = \left( -\frac{t}{y} + \frac{1-t}{1-y} \right) \cdot y(1-y) \cdot x_j$$
$$= (y - t)x_j$$

(verify this)

Gradient descent (coordinatewise) update to find the weights of logistic regression:

$$w_j \leftarrow w_j - \alpha \frac{\partial \mathcal{J}}{\partial w_j}$$
$$= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)}$$

## Multiclass Classification and Softmax Regression

### One-hot Encoding

Targets form a discrete set  $\{1, \dots, K\}$ .

It's often more convenient to represent them as [one-hot vectors](#), or a [one-of-K encoding](#):

$$\mathbf{t} = \underbrace{(0, \dots, 0, 1, 0, \dots, 0)}_{\text{entry } k \text{ is } 1} \in \mathbb{R}^K$$

### Multiclass Linear Classification

分类任务经常不止一类，比如手写数字分类有 10 类。

We can start with a linear function of the inputs.

Now there are  $D$  input dimensions and  $K$  output dimensions, so we need  $K \times D$  weights, which we arrange as a [weight matrix](#)  $\mathbf{W}$ .

Also, we have a  $K$ -dimensional vector  $\mathbf{b}$  of biases.

注：X 为 D 维矩阵，W 为  $K \times D$  维矩阵，b 为 K 维矩阵，这样  $\mathbf{W}\mathbf{X} + \mathbf{b}$  的结果为 K 维矩阵。

A linear function of the inputs:

$$z_k = \sum_{j=1}^D w_{kj} x_j + b_k \quad \text{for } k = 1, 2, \dots, K$$

$$y_i = \begin{cases} 1 & i = \arg \max_k z_k \\ 0 & \text{otherwise} \end{cases}$$

注：上式和 logistic classification 不一样。logistic classification 是之间算出  $\mathbf{W}\mathbf{x}$  然后进行分类。而这里是先对每一类进行预测， $z_k$  的大小可以解释为模型倾向于将  $k$  预测结果的程度。

举个例子，假如做手写数字，一共有 10 类，输入 1 个有  $D$  个 feature 的 input。那么  $\mathbf{w}$  就是  $10 \times D$ ，对于每一类都有特殊的  $D$  个权重。然后模型要对 input 进行 10 次预测，取其中最大的作为结果（在 one-hot encoding 中将相应的位置设为 1，即上面第二个式子所做的）。

We can eliminate the bias  $\mathbf{b}$  by taking  $\mathbf{W} \in \mathbb{R}^{K \times (D+1)}$  and adding a dummy variable  $x_0 = 1$ . So, vectorized:

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad \text{or with dummy } x_0 = 1 \quad \mathbf{z} = \mathbf{W}\mathbf{x}$$

## Softmax Regression

We want soft predictions that are like probabilities, i.e.,  $0 \leq y_k \leq 1$  and  $\sum_k y_k = 1$ .

A natural activation function to use is the [softmax function](#), a multivariable generalization of the logistic function:

$$y_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$$

- Outputs can be interpreted as probabilities (positive and sum to 1)
- If  $z_k$  is much larger than the others, then  $\text{softmax}(\mathbf{z})_k \approx 1$  and it behaves like argmax.

注：上式中， $z_k$  代表  $\mathbf{z}$  中第  $k$  个元素， $\sum_{k'} e^{z_{k'}}$  代表对  $\mathbf{z}$  中每一个元素进行  $e^{z_{k'}}$  后的总和。



If a model outputs a vector of class probabilities, we can use cross-entropy as the loss function:

$$\begin{aligned}\mathcal{L}_{\text{CE}}(\mathbf{y}, \mathbf{t}) &= - \sum_{k=1}^K t_k \log y_k \\ &= -\mathbf{t}^\top (\log \mathbf{y}),\end{aligned}$$

where the log is applied elementwise.

上式叫 softmax-cross-entropy function。

## Gradient Descent

Softmax regression (with dummy  $x_0 = 1$ ):


$$\begin{aligned}\mathbf{z} &= \mathbf{W}\mathbf{x} \\ \mathbf{y} &= \text{softmax}(\mathbf{z}) \\ \mathcal{L}_{\text{CE}} &= -\mathbf{t}^\top (\log \mathbf{y})\end{aligned}$$

Gradient descent updates can be derived for each row of  $\mathbf{W}$ :

$$\begin{aligned}\frac{\partial \mathcal{L}_{\text{CE}}}{\partial \mathbf{w}_k} &= \frac{\partial \mathcal{L}_{\text{CE}}}{\partial z_k} \cdot \frac{\partial z_k}{\partial \mathbf{w}_k} = (y_k - t_k) \cdot \mathbf{x} \\ \mathbf{w}_k &\leftarrow \mathbf{w}_k - \alpha \frac{1}{N} \sum_{i=1}^N (y_k^{(i)} - t_k^{(i)}) \mathbf{x}^{(i)}\end{aligned}$$

## Limits of Linear Classification


Some datasets are not linearly separable, e.g. **XOR**



## Showing that XOR is not linearly separable

通过 contradiction (反证法) 来证明:

- If two points lie in a half-space, line segment connecting them also lie in the same halfspace.
- Suppose there were some feasible weights (hypothesis). If the positive examples are in the positive half-space, then the green line segment must be as well.
- Similarly, the red line segment must lie within the negative half-space.



- But the intersection can't lie in both half-spaces. Contradiction!

## Overcome

- Sometimes we can overcome this limitation using [feature maps](#), just like for linear regression. E.g., for **XOR**:

$$\psi(\mathbf{x}) = \begin{pmatrix} x_1 \\ x_2 \\ x_1 x_2 \end{pmatrix}$$

$x_1$	$x_2$	$\psi_1(\mathbf{x})$	$\psi_2(\mathbf{x})$	$\psi_3(\mathbf{x})$	$t$
0	0	0	0	0	0
0	1	0	1	0	1
1	0	1	0	0	1
1	1	1	1	1	0

- This is linearly separable. (Try it!)